

NetCoffee: a fast and accurate global alignment approach to identify functionally conserved proteins in multiple networks

Jialu Hu^{1,*}, Birte Kehr^{1,2} and Knut Reinert¹¹Department of Mathematics and Computer Science, Freie Universität Berlin, Takustrasse 9, 14195 Berlin, Germany and²Max Planck Institute for Molecular Genetics, Ihnestr. 63-73, 14195 Berlin, Germany

Associate Editor: Anna Tramontano

ABSTRACT

Motivation: Owing to recent advancements in high-throughput technologies, protein–protein interaction networks of more and more species become available in public databases. The question of how to identify functionally conserved proteins across species attracts a lot of attention in computational biology. *Network alignments* provide a systematic way to solve this problem. However, most existing alignment tools encounter limitations in tackling this problem. Therefore, the demand for faster and more efficient alignment tools is growing.

Results: We present a fast and accurate algorithm, *NetCoffee*, which allows to find a global alignment of multiple protein–protein interaction networks. *NetCoffee* searches for a global alignment by maximizing a target function using *simulated annealing* on a set of weighted bipartite graphs that are constructed using a triplet approach similar to *T-Coffee*. To assess its performance, *NetCoffee* was applied to four real datasets. Our results suggest that *NetCoffee* remedies several limitations of previous algorithms, outperforms all existing alignment tools in terms of speed and nevertheless identifies biologically meaningful alignments.

Availability: The source code and data are freely available for download under the GNU GPL v3 license at <https://code.google.com/p/netcoffee/>.

Contact: Jialu.Hu@fu-berlin.de

Supplementary information: Supplementary data are available at *Bioinformatics* online.

Received on May 27, 2013; revised on November 8, 2013; accepted on December 7, 2013

1 INTRODUCTION

Discovering functionally conserved proteins across different species is a fundamental task in comparative systems biology (Park *et al.*, 2011; Tatusov *et al.*, 1997). With the development of high-throughput technologies such as mass spectrometry (Ho, 2002), microarrays (Lashkari *et al.*, 1997), yeast two-hybrid assays (Ito *et al.*, 2001) and next-generation sequencing, a tremendous amount of genomic, proteomics and protein interaction data has been generated and became available in public databases (Szkarczyk *et al.*, 2011; Uniprot Consortium, 2007). This comprehensive experimental data provide a basis for analyses that aim at discovering conservation of protein function among different species, i.e. *functional orthologs* (FO). *Orthologs* are genes/proteins of different species that have descended from the same

gene in a common ancestor. *The FO* are orthologs whose functions are conserved in different species.

It is often assumed that two proteins with similar sequences or similar structure have similar functions, and conversely that functionally related proteins have similar sequences. Based on this assumption, a number of approaches that use sequence similarity have been developed, e.g. reciprocal-best-BLAST-hits (RBH), for predicting FO. This resulted in several orthologs databases such as the *Clusters of Orthologous Groups* (Tatusov *et al.*, 2000), *Inparanoid* (O'Brien *et al.*, 2005) and *OrthoDB* (Waterhouse *et al.*, 2011). However, high sequence similarity does not necessarily indicate functional conservation. Because functional sites of proteins are usually only one or several small parts of the whole sequence, two proteins can have a highly significant overall similarity even though all functional sites are completely different (Brutlag, 2008). To overcome this problem, *network alignment* approaches (Bandyopadhyay *et al.*, 2006; Liao *et al.*, 2009; Shih and Parthasarathy, 2012) have been proposed that supplement sequence-based algorithms with information from protein–protein interaction (PPI) networks.

Network alignment approaches can be generally classified into pairwise or multiple and into local or global approaches. *Pairwise* approaches align two networks and *multiple* approaches three and more networks. *Local* alignment approaches detect conserved subnetworks of two (pairwise local alignment) or more (multiple local alignment) networks. These conserved subnetworks are usually independent high-scoring local regions of the compared networks, each implying a putative functional module such as a protein complex (Sharan, 2005) or a metabolic pathway (Kelley, 2003; Kelley *et al.*, 2004). *Global* alignment approaches determine an overall alignment between the input networks (Singh *et al.*, 2008). The resulting alignments can be used to transfer annotation from characterized to uncharacterized proteins. Typically, pairwise global alignment algorithms attempt to provide only a one-to-one node mapping between proteins of compared networks (El-Kebir *et al.*, 2011; Kuchaiev and Prulj, 2011), whereas multiple global alignment algorithms attempt to provide a many-to-many node mapping (Flannick *et al.*, 2009; Liao *et al.*, 2009). A many-to-many node mapping allows to find a set of functionally similar proteins that have descended from the same protein in a common ancestor based on four types of evolutionary events: protein deletion, protein duplication, protein mutation and paralog mutation (Flannick *et al.*, 2009). For more information on the difference between local and global alignment, see the Supplementary Material.

*To whom correspondence should be addressed.

Many pairwise local and pairwise global alignment tools have been developed in the last decade. Among pairwise local alignment tools are *PathBlast* (Kelley *et al.*, 2004), *MaWISH* (Koyutürk *et al.*, 2005), *NetworkBlast* (Kalaev *et al.*, 2008), *NetAligner* (Pache *et al.*, 2012), *PINALOG* (Phan and Sternberg, 2012) and *SIPINAL* (Aladag and Erten, 2013). A review by Sharan *et al.* (Sharan and Ideker, 2006) extensively discusses some of these methods. Well-known pairwise global alignment tools are *IsoRank* (Singh *et al.*, 2007), *PISwap* (Chindelevitch *et al.*, 2010), *MI-GRAAL* (Kuchaiev and Prulj, 2011), *Natalie* 2.0 (El-Kebir *et al.*, 2011) and *GHOST* (Patro and Kingsford, 2012).

With the increasing availability of PPI networks, the demand for local and global alignment tools of multiple networks has risen. Several multiple alignment tools have been developed, most notably the multiple global alignment tools *Graemlin* 2.0 (Flannick *et al.*, 2009), *IsoRank-N* (Liao *et al.*, 2009) and *SMETANA* (Sahraeian and Yoon, 2013), and the multiple local alignment tool *NetworkBlast-M* (Kalaev *et al.*, 2009). However, these tools have some limitations. *Graemlin* 2.0 requires a training dataset of known alignments to learn its many network-dependent parameters and a phylogenetic tree, which means it can not be applied to species without known alignments or without a phylogenetic tree. *NetworkBlast-M* does not work on networks containing protein nodes with large vertex degree, such as the yeast network in our test datasets. Both *IsoRank-N* and *SMETANA* need a lot of computing time for aligning six or more species. Thus, there is a demand for tools that can deal with many networks in a more efficient way.

As a remedy for these limitations, we present a fast and accurate tool *NetCoffee*, which addresses the problem of global alignment of multiple networks. The algorithm implemented in *NetCoffee* has four main steps: (i) building the PPI networks and a library of bipartite graphs, (ii) assigning an integrated weight to each edge in the bipartite graphs using a triplet extension approach similar to *T-Coffee*, (iii) building a search space that consists of candidate edges (protein pairs) and (iv) simulated annealing (SA) with a large number of iterations of a Metropolis Scheme to maximize a scoring function for global alignments. We ran *NetCoffee* on four datasets consisting of up to six PPI networks. Our results show that *NetCoffee* overcomes the limitations of existing algorithms, outperforms all existing alignment tools in terms of speed and nevertheless identifies a biologically meaningful alignment.

2 METHODS

2.1 Definitions and notation

Let $\{G_1, G_2, \dots, G_k\}$ represent a set of $k \geq 3$ PPI networks. Each network $G_i = (V_i, E_i)$ is an unweighted graph, where V_i is a set of nodes representing proteins and E_i a set of binary interactions appearing in the networks. We refer to elements of E_i as *interactions* to distinguish them from edges in a different type of graph below. Let $V = \bigcup_{i=1}^k V_i$ be the union of all proteins. A *match set* ϑ is a subset of V . By definition, a *global alignment* of the k networks is a node mapping that consists of a set of mutually disjoint match sets, $\{\vartheta^1, \vartheta^2, \dots, \vartheta^m\}$ with $\vartheta^i \cap \vartheta^j = \emptyset$, $\forall i, j, i \neq j$. A match set can contain more than one node from each network.

Like a match set, a *k-spine* is a subset of V but contains exactly one protein from each network. In addition, two different k-spines can share

nodes. In contrast to global alignments, a *local alignment* of k networks is a set of independent high-scoring local node mappings, each node mapping consisting of a set of *k-spines* (Kalaev *et al.*, 2009).

2.2 Generating a bipartite graph library

Given k species and their corresponding PPI networks, we build a bipartite graph library, which contains a graph $B_{ij} = (V_i \cup V_j, E_{ij})$ for each pair of input networks G_i and G_j , $i \leq j$, $i, j \in \{1, 2, \dots, k\}$. We use the term *edges* to refer to elements in E_{ij} . To determine the sets E_{ij} , we perform an all-against-all sequence comparison with the program BLASTP (Altschul *et al.*, 1997) for each pair of species, including pairs of the same species like human-human. Then, the set of $\binom{k+1}{2}$ bipartite graphs can be constructed by simply joining protein pairs $v_1 \in V_i$, $v_2 \in V_j$ that have an *e-value* $\leq 10^{-7}$ by edges $(v_1, v_2) \in E_{ij}$. In bipartite graphs B_{ii} of the same species, we add only edges for pairs of two distinct proteins $v_1 \neq v_2$ to E_{ii} . This allows us to construct match sets that might reflect duplication events within a species and hence exhibit functional relation within a species.

2.3 Integration of two conservation measures

To search for a biologically meaningful alignment, we developed a linear scoring model that assigns a weight to each edge of the bipartite graphs. The development of the scoring model was intuitively guided by two basic assumptions: (i) functionally conserved proteins are likely to have sequence similarity and (ii) interactions among orthologous proteins are likely to be conserved across species. Likewise, our scoring model consists of two independent parts for sequence and topology similarity. Given an edge $e = (v_1, v_2)$, we use $S_r(v_1, v_2)$ to denote a normalized sequence score and $S_t(v_1, v_2)$ to denote a normalized topology score for proteins v_1 and v_2 . A combined score for the edge e is calculated with $S(v_1, v_2) = \alpha S_r(v_1, v_2) + (1 - \alpha) S_t(v_1, v_2)$, where α is a user-defined parameter controlling how much of the topology score contributes to $S(v_1, v_2)$.

To compute the sequence-based score $S_r(v_1, v_2)$ for a pair of proteins v_1 and v_2 , we adopt a previously introduced log-ratio scoring function that uses distributions of *e-values* in two models, the homology model H and the null model N (Flannick *et al.*, 2006). The null model includes all pairs of proteins from the input networks, whereas the homology model includes only pairs of proteins with an *e-value* $\leq 10^{-7}$. Given the distributions of *e-values* in these two models, we calculate the probabilities to observe the *e-value* $x_{v_1 v_2}$ of the two proteins v_1 and v_2 in the two models, $Pr(x_{v_1 v_2} | H)$ and $Pr(x_{v_1 v_2} | N)$. Our normalized sequence score is the log-ratio

$$y_{v_1 v_2} = \log \frac{Pr(x_{v_1 v_2} | H)}{Pr(x_{v_1 v_2} | N)}$$

of these probabilities scaled to the range from 0 to 1 with the minimal observed log-ratio y_{\min} and maximal observed log-ratio y_{\max} of all protein pairs in the H model:

$$S_r(v_1, v_2) = \frac{y_{v_1 v_2} - y_{\min}}{y_{\max} - y_{\min}}.$$

To compute the topology-based score $S_t(v_1, v_2)$, we use a triplet approach that bears similarities to the concept of overlapping weights (Morgenstern, 1999) and T-Coffee's consistency approach (Notredame *et al.*, 2000) in multiple sequence alignment. Our approach is an incremental process with the final score reflecting the likelihood of a pair of proteins being topologically conserved. Initially, we set the topology-based scores of all edges in the $\binom{k}{2}$ bipartite graphs of two different species to zero. After this initialization, each of the edges has an equal right to be part of the global alignment with regard to the topology similarity. Figure 1a illustrates an example of species that are numbered 1 and 2. Next, we do a series of triplet comparisons as displayed in Figure 1b-e. A *triplet* is a set of three PPI networks and the three involved bipartite graphs. We can construct a series of triplets by combining any

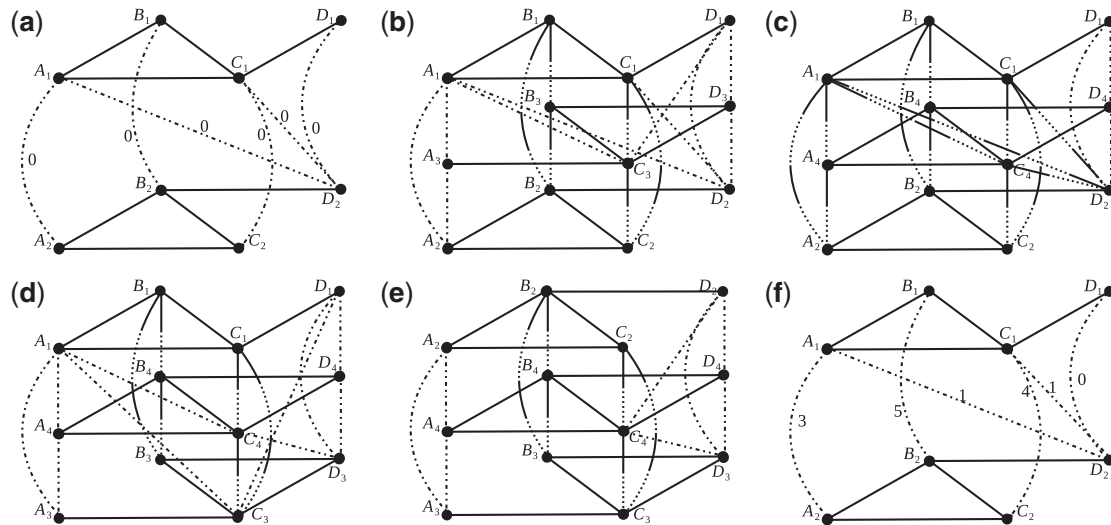


Fig. 1. The workflow of our triplet approach on an example with four species. Proteins are represented by nodes, PPIs by solid lines and edges of bipartite graphs by dashed lines. (a) Initialization of the bipartite graph for species 1 and 2. (b–e) Comparison of triplets. Increasing edge scores of pairs of triplet matches whose proteins share three PPIs, e.g. the edges in line with fine dots. (f) The final topology scores of edges between species 1 and 2

three different PPI networks. A set of three nodes that are mutually connected by edges is a *triplet match*, e.g. $\{A_1, A_2, A_3\}$ in Figure 1b. In the process of reweighing, we consider all pairs of triplet matches that are connected by conserved interactions in all three networks, such as the edges in line with fine dots in Figure 1b–e, and increase the score of each edge in the pair of triplet matches by one. In graphs of the same species, we set the topology-based score of all edges to zero.

All edge scores of the two species 1 and 2 are illustrated in Figure 1f. As an example, the overall topology-based score for the two proteins B_1 and B_2 in Figure 1f is five, which is explained as follows: In Figure 1b the conserved interaction between $\{B_1, C_1\}$ and $\{B_2, C_2\}$ is confirmed by $\{B_3, C_3\}$, and hence the triplet matches $\{B_1, B_2, B_3\}$ and $\{C_1, C_2, C_3\}$ are completely connected by interaction edges contributing one to the score. In Figure 1b, the triplet matches $\{A_1, A_2, A_3\}$ and $\{B_1, B_2, B_3\}$ do not contribute because of the missing interaction edge $\{A_3, B_3\}$. In Figure 1c, the four combinations of triplet matches $\{\{B_1, B_2, B_4\}, \{A_1, A_2, A_4\}\}$, $\{\{B_1, B_2, B_4\}, \{C_1, C_2, C_4\}\}$, $\{\{B_1, B_2, B_4\}, \{A_1, D_2, C_4\}\}$ and $\{\{B_1, B_2, B_4\}, \{C_1, D_2, C_4\}\}$ contribute four to the score. For more details for this example including the other weighted bipartite graphs and pseudo-code for the triplet comparison see the Supplementary Material.

After this process, each edge of the bipartite graphs has been assigned a topology-based score, which we normalize to the range between 0 and 1. However, the distribution of these scores is extremely non-uniform, as the connectivity of biological networks follows a power-law distribution (Barabasi and Albert, 1999). A few edges between hub-nodes have a topology score close to the maximal score, and many others are close to zero. For example, $\sim 90\%$ of the protein pairs have a normalized topology score between 0 and 0.1 in our dataset 2. In contrast, 95% of the protein pairs has a normalized sequence score between 0.6 and 1. This implies that small normalized topology scores might still be statistically significant and indicate a high probability of functional relatedness. A large number of protein pairs have a small normalized topology score because the maximal score is large. Therefore, we lift the small scores up using a power-law redistribution to make sure that the topology score has a reasonable impact on the whole alignment score (for details see the Supplementary Material). This concludes the computation of the edge scores $S(v_1, v_2)$, where each score now reflects sequence similarity and topology conservation.

2.4 Alignment algorithm

Our algorithm for aligning multiple networks first collects candidate edges from the $\binom{k+1}{2}$ bipartite graphs. Subsequently, the algorithm combines some of these candidate edges to a global alignment with the meta-heuristic method SA (Kirkpatrick *et al.*, 1983). The collection of candidate edges reduces the computational complexity while retaining the sensitivity and specificity of the algorithm in praxis.

2.4.1 Collection of candidate edges We have given $\binom{k+1}{2}$ weighted bipartite graphs, $\binom{k}{2}$ of which formed by proteins from two different species. The weights of all edges in $B_{ij}, i < j$ reflect the likelihoods of the edges to be a true match of the global alignment, including information about sequence and topology conservation. We use a maximum weighted matching algorithm, namely, Edmond's Algorithm (Galil, 1983), to find a one-to-one node mapping in each of the $\binom{k}{2}$ bipartite graphs and collect the matching edges as candidate edges. Furthermore, we collect protein pairs of the same species with scores higher than a threshold $\sigma = \eta(1 - \alpha)$. The parameter η is user-defined and enables our method to identify match sets formed by proteins of one species. The term $(1 - \alpha)$ accounts for the fact that the topology score of these edges is always 0. We obtain a collection of candidate edges, denoted as Ω .

2.4.2 Multiple alignment To find a multiple global alignment $\mathbb{A} \subseteq \Omega$, we define the scoring function $\Phi(\mathbb{A}) = \sum_{\vartheta \in \mathbb{A}} f(\vartheta)$, where $f(\vartheta)$ is the score of a match set $\vartheta = \{v_1, v_2, \dots, v_{|\vartheta|}\}$. The score of ϑ is calculated with the function $f(\vartheta) = \sum_{i,j} S(v_i, v_j) \delta_{ij}$, where $\delta_{ij} = 1$ if $\{v_i, v_j\} \in \Omega$, otherwise $\delta_{ij} = 0$.

Let I be the collection of all possible global alignments. Then, the problem of multiple global alignment can be modeled as an optimization problem $\max_{\mathbb{A} \in I} \Phi(\mathbb{A})$. We use an SA approach to approximate the highest-scoring alignment. Unlike the strategy of progressive alignment (Flannick *et al.*, 2006), which successively aligns closest pairs of networks and constructs a new network alignment, our SA approach starts with an empty alignment of all networks and runs a large number of iterations of a *Metropolis* scheme (Metropolis *et al.*, 1953) to maximize $\Phi(\mathbb{A})$.

Algorithm 1 Simulated annealing algorithm

Input: Matching edges Ω , K , T_{\min} , T_{\max} , s
Output: A solution \mathbf{x}^* with a set of mutually disjoint match sets

```

1:  $\mathbf{x} = \emptyset$ ,  $T_0 = T_{\max}$ ,  $i = 1$ ;
2: while  $i \leq K$  do
3:    $n = 0$ ;
4:    $T_i = T_0 - \frac{i(T_{\max} - T_{\min})}{K}$ ;
5:   while  $n < N$  do
6:     draw arbitrary sample  $\xi \in \Omega$  from uniform distribution;
7:      $\mathbf{x}' = \text{updateState}(\mathbf{x}, \xi)$ ;
8:      $\Delta\Phi = \Phi(\mathbf{x}') - \Phi(\mathbf{x})$ ;
9:     if  $\Delta\Phi > 0$  then
10:       $\mathbf{x} = \mathbf{x}'$ ;
11:     else  $\text{rand}(0, 1) < \exp\{\Delta\Phi/(sT_i)\}$ 
12:       $\mathbf{x} = \mathbf{x}'$ ;
13:     end if
14:      $n = n + 1$ ;
15:   end while
16:    $i = i + 1$ ;
17: end while
18:  $\mathbf{x}^* = \mathbf{x}$ ;
19: return  $\mathbf{x}^*$ ;
```

The pseudo-code of the SA approach is given in Algorithm 1. Let $\mathbf{x} \in I$ be a feasible solution (a set of mutually disjoint match-sets) for the problem and $\Phi(\mathbf{x})$ the alignment score of \mathbf{x} . At the beginning of the algorithm, we initialize our alignment \mathbf{x} with \emptyset and set a temperature parameter T_0 to its maximum. In the following annealing phase, we decrease the temperature and repeatedly perturb the current solution \mathbf{x} with a *Metropolis* scheme using $\pi_i \propto \exp(-\Phi(\mathbf{x})/(sT_i))$ as the equilibrium distribution. Parameters s, K, N, T_{\min} and T_{\max} control the SA. The *updateState* (\mathbf{x}, ξ) updates the current alignment with an arbitrary sample $\xi = \{u, v\} \in \Omega$. It runs into four possible scenarios. Let $u \notin \mathbf{x}$ indicate that $\forall \zeta \in \mathbf{x}$, $u \notin \zeta$, and $u \in \mathbf{x}$ indicate that $\exists \zeta \in \mathbf{x}$, such that $u \in \zeta$. Then, the scenarios are (i) $u \notin \mathbf{x}$ and $v \notin \mathbf{x}$; (ii) $u \notin \mathbf{x}$ and $v \in \mathbf{x}$; (iii) $u \in \mathbf{x}$ and $v \notin \mathbf{x}$; and (iv) $u \in \mathbf{x}$ and $v \in \mathbf{x}$, but u and v are not in the same match set. In the first scenario, u and v are added to the current alignment \mathbf{x} as a new match set. In the other scenarios, u and v are moved to the same match set of the alignment \mathbf{x} in two possible ways, called *combination* and *substitute*. Details are described in the Supplementary Material. We continue this process until the ‘temperature’ T_i decreases to T_{\min} .

2.5 Complexity analysis

We assume that the number of proteins in the largest PPI network is n , and the number of input networks is k . The pseudo-code of the triplet approach (see Supplementary Material) has a complexity of $\binom{k}{3} O(n^6)$. Suppose there is a bipartite graph, $B_s = (V_{s1} \cup V_{s2}, E_s)$, the running time complexity of *Edmond’s Algorithm* on B_s is $O(|V_{s1} \cup V_{s2}| \cdot \log |E_s|)$. Therefore, the collection of candidate edges costs $\binom{k}{3} O(n \log(n))$ time.

The convergence time of SA has been a widely studied question in the last two decades. We assume $\Delta = \max\{\Phi(\mathbf{x}') - \Phi(\mathbf{x})\}$, where \mathbf{x}' is a neighbor state of state \mathbf{x} . As shown by the proof in Rajasekaran (1990), SA converges (at any temperature) in time $2\beta[d \exp(\Delta/(sT))]^D$, where D is the diameter, d is the degree of the underlying Markov chain and β is defined by the convergence probability $\geq (1 - 2^{-\beta})$. Theoretically, D and d are hard to calculate. However, in practice, the complexity of SA only depends on two parameters of the cooling scheme, K and N . From Algorithm 1, we can easily find that the complexity is $\Omega(K \cdot N)$, which is independent of the number of compared species k . To sum up, practically, *NetCoffee* is able to deal with multiple networks and has a favorable time complexity. Our results show that the alignment score converges rapidly in our experiments (see Supplementary Fig. S5).

3 RESULTS AND DISCUSSION

3.1 Test datasets

We have evaluated *NetCoffee* on three datasets of up to five eukaryotic species and one dataset of six microbes as shown in Table 1. The five eukaryotic species include *Homo sapiens* (human), *Mus musculus* (mouse), *Drosophila melanogaster* (fruit fly), *Caenorhabditis elegans* (nematode) and *Saccharomyces cerevisiae* (yeast). The six microbes include *Escherichia coli*, *Salmonella typhimurium*, *Vibrio cholerae*, *Campylobacter jejuni* NCTC 11168, *Helicobacter pylori* 26695 and *Caulobacter crescentus*.

To build the five eukaryotic networks of dataset 0, 1 and 2, we collected all experimentally determined interactions from the public database IntAct (Kerrien *et al.*, 2012). In addition, we collected the reference proteome sets of the five species from UniProtKB/Swiss-Prot release 2012_07 (UniProt Consortium, 2007), which are used for all-against-all sequence comparisons. To make sure the proteins in our networks are non-redundant and well-annotated, we discarded interactions between proteins that are not in the reference proteome sets. The number of proteins and interactions of these PPI networks are given in Table 1. Dataset 3 is the same dataset used in the original publication of *Graemlin* 2.0 (Flannick *et al.*, 2009).

For analyzing the biological quality of the alignments, gene ontology (GO) information was collected from UniProt-GOA (Camon *et al.*, 2004) (downloaded on Jan. 8, 2013) to annotate proteins with the three basic types of ontologies: *biological process* (BP), *molecular function* (MF) and *cellular component* (CC). To exclude unreliable function annotations, GO annotations with evidence codes IEA (inferred from electronic annotation) and ISS (inferred from sequence or structural similarity) were discarded.

3.2 Experimental setup

We have implemented *NetCoffee* in C++ using the *LEMON Graph Library* (Dezs *et al.*, 2011) version 1.2.3. The implementation supports multicore parallelism for the triplet comparison. We ran *NetCoffee* on all four datasets and tuned its SA parameters such that the SA process converges to a stable score (see Supplementary Fig. S5). The default values are now $s = 0.005$, $K = 100$, $N = 2000$, $T_{\min} = 10$, $T_{\max} = 100$ and $\eta = 1.0$.

To compare *NetCoffee* with the state-of-the-art algorithm *IsoRank-N*, we executed *IsoRank-N* on the same datasets with recommended parameters: $K = 20$, $\text{thresh} = 10^{-4}$, $\text{maxveclen} = 10^6$. Additionally, *NetworkBlast-M*, *Graemlin* 2.0 and *SMETANA* were included in our assessment. However, *NetworkBlast-M* is unable to work on dataset 0, 2 and 3 for two reasons. First, the yeast network has proteins with up to 3276 interactions, which is prohibitive for *NetworkBlast-M*. Second, *NetworkBlast-M* requires *e*-values as a protein similarity measure, but dataset 3 provides only bitscores. Furthermore, we ran *Graemlin* 2.0 only on dataset 3 because it needs additional training data (i.e. known alignments for the compared species) to learn its parameters. Because *Graemlin* 2.0 identifies match sets whose proteins are from a single species, we set $\eta = 0.7$ for dataset 3 to allow a fair comparison with *Graemlin* 2.0.

Table 1. The number of proteins and PPI of four datasets that consist of the PPI networks from 11 species

Species	Proteins	Interactions	Dataset 0	Dataset 1	Dataset 2	Dataset 3
<i>H.sapiens</i>	8777	28 366		✓	✓	
<i>M.musculus</i>	1531	1626		✓	✓	
<i>D.melanogaster</i>	1534	2664	✓	✓	✓	
<i>C.elegans</i>	767	915	✓	✓	✓	
<i>S.cerevisiae</i>	5739	36 226	✓		✓	
<i>E.coli</i>	4179	169 636				✓
<i>V.cholerae</i>	3044	76 341				✓
<i>C.jejuni</i> 11168	1424	76 913				✓
<i>H.pylori</i> 26695	1206	48 430				✓
<i>C.crescentus</i>	3022	52 302				✓
<i>S.typhimurium</i>	4326	151 118				✓

We input the networks of the species in the same order for all programs, namely, the order from Table 1. Only the results of *IsoRank-N* depend on the order of input species. All experiments mentioned in the following parts were carried out on the same machine, an Intel(R) Xeon(R) CPU X5550 with 2.67GHz.

3.3 Performance comparison

We demonstrate the quality of our alignments in terms of coverage and consistency and assess the performance of our method by measuring running times. Coverage, which serves as a proxy for sensitivity, indicates the amount of input data the algorithm can explain. Consistency, which serves as a proxy for specificity, measures the functional similarity of proteins in each match set. Coverage can be easily achieved by sacrificing consistency and vice versa. The running time demonstrates the ability of *NetCoffee* to deal with large datasets. Intuitively, the goal is to find a global alignment that has a good consistency while explaining as many proteins as possible (i.e. high coverage) in reasonably short time. We first look at differences the programs exhibit in coverage and then investigate the consistency of the match sets with three measures. Next, we compare running times and, finally, demonstrate how much *NetCoffee* benefits from the integration of similarity and topology score by addressing the influence of the parameter α .

3.3.1 Coverage For each program, we calculated the percentage of proteins (PPV) in the whole set of proteins that are covered by the alignment as the coverage (see Table 2). In comparison with *IsoRank-N* and *NetworkBlast-M*, the coverage of *NetCoffee* is significantly higher. For instance, the PPV of *NetCoffee* is up to 41.8% for dataset 1, whereas it is only 31.1% for *IsoRank-N* and 16.1% for *NetworkBlast-M*. The lower coverage of these two alignment tools can be explained by the facts that *NetworkBlast-M* is a local aligner and, thus, considers only conserved modules; *IsoRank-N* aligns proteins of at least three species into match sets and does not report match sets of proteins from only two species (These match sets can be recognized by running the pairwise aligner *IsoRank* on each pair of species.) (see an example in Supplementary Table S1). In comparison with *Graemlin 2.0*, *NetCoffee* also has a slightly higher PPV value except for the extreme case of $\alpha = 1$. When $\alpha = 1$, sequence

scores of all pairs of proteins are set to 0 in *NetCoffee*. As a result, all protein pairs from a single species are excluded from the collection of candidate edges and consequently from the alignment. Hence, the coverage drops to 69.7% for dataset 3. In comparison with *SMETANA*, the coverage of *NetCoffee* is similar. *NetCoffee* achieves a lower PPV for dataset 0, 1 and 2, but a higher PPV for dataset 3. Concerning the number of match sets, *IsoRank-N* identifies more match sets formed by proteins from three of the compared species, and both *Graemlin 2.0* and *SMETANA* find more match sets for dataset 3 than *NetCoffee* except for $\alpha = 1$ (see Supplementary Table S1).

3.3.2 Consistency An alignment tool that achieves a high coverage is not necessarily better than others. For example, a random global alignment may cover all proteins but aligns many unrelated proteins. Hence, we now address the performance of the alignment tools in terms of consistency. Consistency demonstrates the biological significance of predicted match sets.

As a first consistency measure, we computed the mean entropy and the mean normalized entropy of the predicted match sets in the alignments of each algorithm. We calculated the entropy of a match set with the same method as in *IsoRank-N* according to its GO annotations. A match set has lower entropy if its GO annotations are more functionally coherent. From Table 2, we can see that the entropy of *NetCoffee* is considerably lower than that of *IsoRank-N* and *NetworkBlast-M* no matter which α was used, whereas at the same time having a high coverage. Additionally, the entropy of *NetCoffee* is lower than that of *SMETANA* on all datasets except for dataset 3. In comparison with *Graemlin 2.0*, *NetCoffee* achieves nearly identical entropy results for dataset 3, whereas being considerably faster. The results for $\alpha = 0$ and $\alpha = 1$ demonstrate that both of our two conservation measures can favorably predict the functional relatedness between protein pairs.

Dataset 3 exhibits an interesting trade-off using the α parameter in terms of coverage and consistency. For $\alpha = 1$, *NetCoffee* has the lowest entropy, however, at the cost of a much lower coverage. Decreasing α improves the coverage while deteriorating the entropy measure. This behavior is less pronounced for the other datasets. However, it shows that the α parameter can be used for having a specificity versus sensitivity trade-off.

Table 2. Coverage, entropy and speed comparison

Dataset	Measure	NetCoffee							IsoRank-N							NBM	Gr. 2.0	SME
		$\alpha = 0.0$	$\alpha = 0.3$	$\alpha = 0.4$	$\alpha = 0.5$	$\alpha = 0.6$	$\alpha = 0.7$	$\alpha = 1.0$	$\alpha = 0.0$	$\alpha = 0.3$	$\alpha = 0.4$	$\alpha = 0.5$	$\alpha = 0.6$	$\alpha = 0.7$	$\alpha = 1.0$	-	-	-
D-0	PPV (%)	28.3	28.2	28.4	28.3	28.1	28.3	27.6	6.75	16.1	16.1	15.7	15.1	14.2	2.14	*	*	35.2
	ME	1.525	1.504	1.511	1.519	1.522	1.528	1.523	3.562	2.927	2.941	2.952	3.083	3.057	3.235	*	*	2.393
	MNE	0.6081	0.6026	0.6031	0.6054	0.6052	0.6059	0.6091	1.093	0.9627	0.967	0.9705	0.9993	0.9942	0.9951	*	*	0.8374
	Time	2s	2s	2s	2s	2s	2s	2s	9.9m	20.5m	21.7m	25.3m	32.5m	41m	70.5m	*	*	53s
D-1	PPV (%)	41.8	41.2	41	41	40.5	40.2	41	19.4	31.1	30.9	30.4	30.2	29.7	10.8	16.1	*	52.6
	ME	2.603	2.645	2.6	2.615	2.608	2.614	2.589	4.366	3.927	3.896	3.986	3.992	4.043	4.863	4.130	*	3.054
	MNE	0.8642	0.8721	0.8638	0.8669	0.8651	0.8652	0.8601	1.258	1.173	1.167	1.186	1.185	1.194	1.336	1.227	*	0.9616
	Time	22s	21s	21s	21s	22s	21s	22s	14.6m	29.6m	31.4m	37.7m	46.9m	64.2m	101m	13.5m	*	3.3m
D-2	PPV (%)	49.5	49.1	48.8	48.6	48.3	48.1	47.1	21.9	33.8	33.3	32.7	32.1	31.8	9.67	*	*	58.1
	ME	2.257	2.288	2.265	2.286	2.293	2.277	2.238	3.942	3.597	3.629	3.645	3.681	3.686	4.403	*	*	2.592
	MNE	0.7918	0.7988	0.7931	0.7979	0.7995	0.7954	0.7883	1.167	1.103	1.109	1.113	1.12	1.12	1.244	*	*	0.8656
	Time	48.2s	47.3	51.9s	55.6s	46.6s	47.8s	48.0s	1.70h	3.01h	3.77h	4.18h	4.81h	5.86h	9.05h	*	*	6.2m
D-3	PPV (%)	84.8	84.3	84.3	84.1	83.9	83.6	69.7	*	*	*	*	*	*	*	*	82.4	79.3
	ME	0.267	0.266	0.267	0.268	0.266	0.267	0.249	*	*	*	*	*	*	*	*	0.263	0.248
	MNE	0.203	0.203	0.204	0.205	0.204	0.205	0.201	*	*	*	*	*	*	*	*	0.205	0.199
	Time	4.9m	5.9m	5.3m	5.8m	5.7m	7.3m	3.9m	> 72h	>72h	>72h	>72h	>72h	>72h	>72h	*	6.7h	7.8m

Note: The five algorithms *NetCoffee*, *IsoRank-N*, *NetworkBlast-M* (NBM), *Graemlin 2.0* (Gr. 2.0) and *SMETANA* (SME) were tested on the four datasets. The rows list the PPV, mean entropy (ME), mean-normalized entropy (MNE) and the running time. D-x in the first column represents the test dataset, Dataset-x. S, m and h in the row of time represent seconds, minutes and hours. Bold face numbers represent the best performance with respect to each row. The parameter of α in both *NetCoffee* and *IsoRank-N* demonstrates the percentage of the topology score contributing to the whole-alignment score. And '-' indicates α is not a parameter of the corresponding aligner, '*' means the corresponding aligner is not applicable in this dataset.

Second, we assessed consistency by three elaborate semantic similarity measures introduced in Schlicker *et al.* (2006, 2007): *BPscore*, *MFscore* and *rfunSim*. Unlike many existing approaches (El-Kebir *et al.*, 2011; Kuchaiev and Prulj, 2011) that simply evaluate functional similarity by counting the number of common GO terms of involved proteins, *BPscore* and *MFscore* assess the functional similarity of two proteins by exploiting BP and MF annotations with the GO hierarchy tree. The measure *rfunSim* is a combination of *BPscore* and *MFscore* (for details see the Supplementary Material). We report the arithmetic mean of the similarity scores of all involved protein pairs as the functional consistency of a match set. For instance, given a match set $\vartheta = (v_1, v_2, \dots, v_{|\vartheta|})$, the functional consistency of ϑ with respect to the BP annotation is defined as

$$\overline{BPscore}(\vartheta) = \frac{\sum_{i \neq j} BPscore(v_i, v_j)}{\binom{|\vartheta|}{2}}, \quad i, j \in \{1, 2, \dots, |\vartheta|\}.$$

Analogously, we can calculate $\overline{MFscore}$ and $\overline{rfunSim}$. All three scores range from 0 to 1, which translates into an increasing degree of functional similarity. We calculated the scores using the functional similarity search tool (FSST) (Schlicker *et al.*, 2007). To avoid skipping too many meaningful match sets, match sets that contain <40% uncharacterized proteins were also taken into consideration. We separately compared match sets that contain proteins from 3, 4 and 5 species. And the distribution of match sets in each category can be seen in Supplementary Table S1.

We compared the consistency of *NetCoffee* with that of *IsoRank-N* (see Fig. 2) and *SMETANA* (see Supplementary

Fig. S6) on their alignments of dataset 2. As shown in Figure 2a–c, when $\alpha > 0$, the $\overline{BPscore}$ of *NetCoffee* is higher than that of *IsoRank-N*, and the $\overline{MFscore}$ and $\overline{rfunSim}$ are roughly the same. More importantly, the advantage of *NetCoffee* expands when i (i.e. the number of species) increases to 4, as shown in Figure 2d–f, although it identifies more match-sets. *NetCoffee* shows significant improvements with regard to the $\overline{BPscore}$, $\overline{MFscore}$ and $\overline{rfunSim}$ except for the case of $\alpha = 0$. When $\alpha = 0$, *IsoRank-N* reaches its highest point. However, we do not recommend to use $\alpha = 0$ for *IsoRank-N*, as its coverage drops to only 21.9%. For $i = 5$ illustrated in Figure 2g–i, *IsoRank-N* improves the quality of match sets in terms of $\overline{BPscore}$. The two algorithms are comparable in terms of $\overline{MFscore}$ and $\overline{rfunSim}$. However, *NetCoffee* identifies ~3–8 times more match sets than *IsoRank-N* (see Supplementary Table S1). Compared with the alignment of *SMETANA*, match sets identified by *NetCoffee* have lower semantic scores for $i = 3$ but roughly the same scores for $i = 4$ and $i = 5$ (see Supplementary Fig. S6).

Finally, we measured the consistency by computing the percentage of qualified match sets the algorithms identified. As demonstrated in Schlicker *et al.* (2006), almost 60% of protein pairs in the Inparanoid Orthologs (IO) dataset has an $\overline{MFscore} > 0.8$ and 65% has a $\overline{BPscore} > 0.6$. Therefore, we regard those match sets that have an $\overline{MFscore} > 0.8$ or a $\overline{BPscore} > 0.6$ as *qualified match sets*, i. e. functionally related proteins. With these thresholds, ~45% of the match sets recognized by *NetCoffee* are qualified match sets (see Supplementary Fig. S7), which is significantly more than those identified by *IsoRank-N* (~25%) and more than those identified by *SMETANA* (~42%). Visualizations of the GO trees for each qualified match set [drawn using the package *GO::TermFinder*

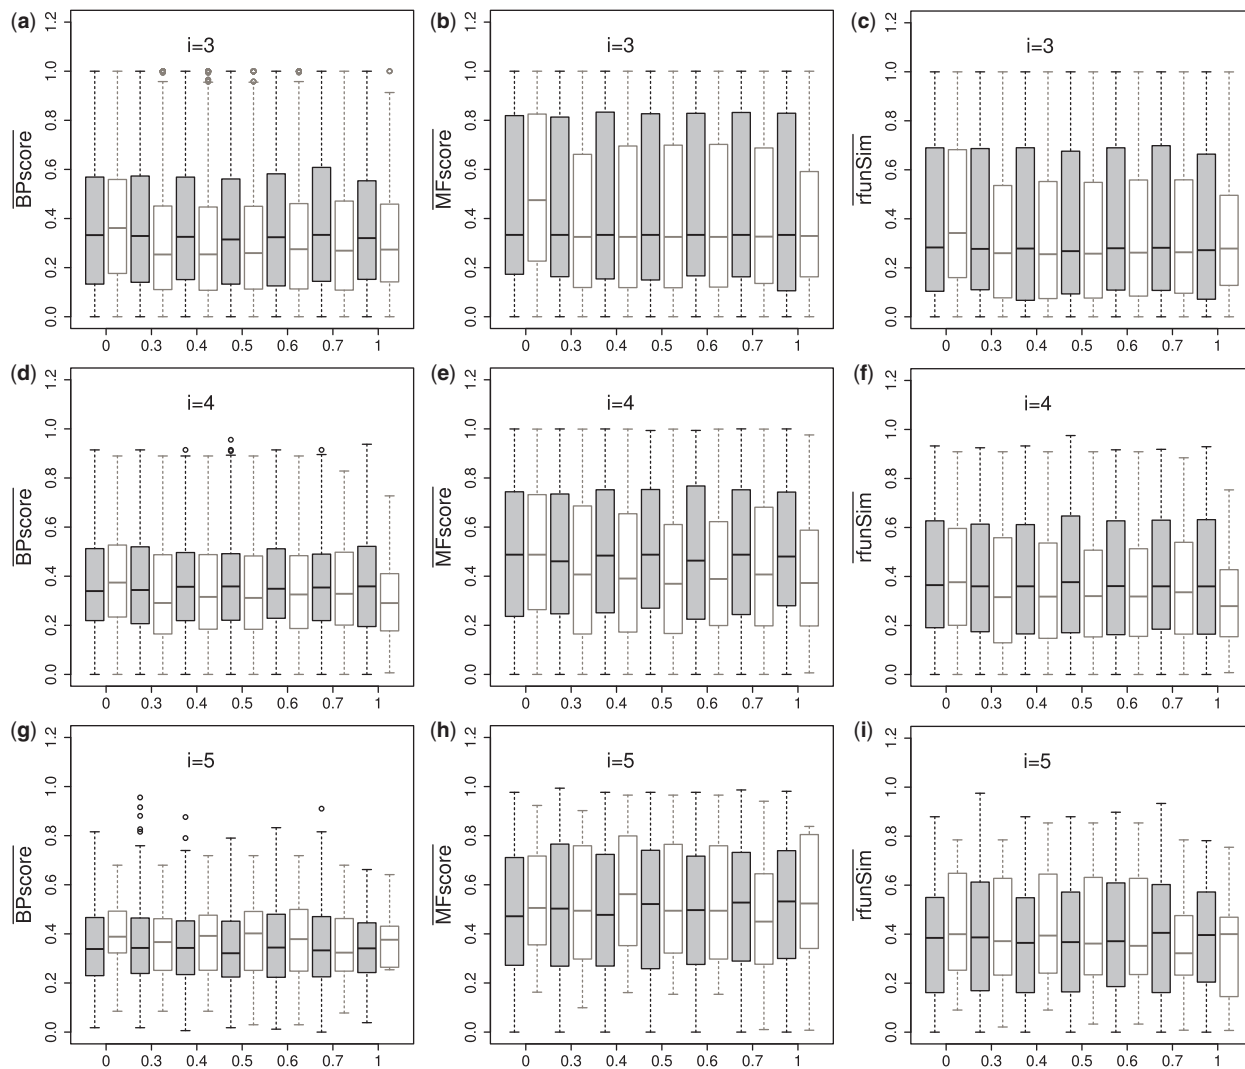


Fig. 2. Consistency comparison on dataset 2 between *NetCoffee* (gray boxes) and *IsoRank-N* (white boxes). Box plots for the semantic similarity measures *BPscore*, *MFscore* and *rfunSim* of match sets conserved by $i \in \{3, 4, 5\}$ species, with respect to the parameter α (the horizontal axis)

(Boyle *et al.*, 2004)] and more information of the alignment with $\alpha = 0.3$ are available for download from <https://code.google.com/p/netcoffee/downloads/list>.

3.3.3 Running time Table 2 demonstrates that our method is robust to the parameter α in terms of running time. The running time of *IsoRank-N*, however, increases dramatically when α grows. Specifically, *NetCoffee* is ~ 1 –3 orders of magnitude faster than *IsoRank-N*, 37 times faster than *NetworkBlast-M*, 82 times faster than *Graemlin 2.0* (including training time) and 2–26 times faster than *SMETANA*. We chose to report the results achieved with multiple cores (i.e. eight cores) because they are the real running time for *NetCoffee*. *NetCoffee* is still faster than its competitors even on a single core except for *SMETANA* (see Supplementary Table S2).

3.3.4 Influence of the parameter α To figure out how much the alignment tools benefit from the topology and sequence score, we ran both *NetCoffee* and *IsoRank-N* with various α values. If

$\alpha = 0$, the global alignment is constructed only based on sequence score, and if $\alpha = 1$, only based on topology score.

Table 2 and Figure 2 demonstrate that *NetCoffee* is robust to the parameter α in terms of coverage, consistency and speed, and that the α parameter can be used for having a specificity versus sensitivity trade-off. Both the topology and the sequence score favorably predict functional relatedness between protein pairs.

However, using either sequence score or topology score alone is not favorable for the coverage of *IsoRank-N* as shown in Table 2. Furthermore, the alignment quality and the computing time depend on α . Table 2 suggests that the performance of *IsoRank-N* tends to be best at a value of $\alpha = 0.3$.

4 CONCLUSION

We have proposed a fast and accurate algorithm for global alignment of multiple networks. It overcomes several limitations of existing tools by aligning multiple networks without additional

training data, finding a global alignment of six species within several minutes and scaling to networks with tens of thousands of proteins and interactions. Further, it is the first alignment tool that can run with multiple cores in parallel.

We rigorously combine protein sequence similarity and network topology similarity into a suitable scoring scheme for multiple networks, adapting a successful technique from multiple sequence alignment. This allows us to model the problem as a combinatorial optimization problem, which we solve with *SA*. On PPI networks of five eukaryotic species, such as human, mouse, fruit fly, nematode and yeast, our implementation *NetCoffee* successfully finds a global alignment covering ~50% of the proteins; and ~45% of the match sets are qualified.

We compared *NetCoffee* with four existing tools, three of which fail to run on at least one of the three test datasets in our benchmark. The results indicate that *NetCoffee* outperforms the state-of-the-art algorithm *IsoRank-N* in terms of coverage and consistency, and at the same time is ~1–3 orders of magnitude faster. Compared with *NetworkBlast-M*, *Graemlin* 2.0 and *SMETANA*, *NetCoffee* not only overcomes their limitations but also retains the quality of alignments in terms of both coverage and consistency.

This suggests that *NetCoffee* provides substantial improvements to global network alignment and that the research community working on function annotation and phylogenetic analysis can benefit from it. Further, its application is not restricted to PPI networks. It could also be extended to other types of complex networks, such as Scientific Collaboration Networks (SCN) and World Wide Web Networks (WWWN).

ACKNOWLEDGEMENT

The authors are grateful to Gunnar Klau and Mohammed El-Kebir for helpful discussions. They thank Roded Sharan and Maxim Kalaev for their helpful support regarding *NetworkBlast-M*. They also wish to thank Jason Flannick and Tony Novak for their help with the dataset used in *Graemlin* 2.0.

Funding: China Scholarship Council (CSC).

Conflict of Interest: none declared.

REFERENCES

- Aladag,A.E. and Erten,C. (2013) Spinal: scalable protein interaction network alignment. *Bioinformatics*, **29**, 917–924.
- Altschul,S.F. *et al.* (1997) Gapped BLAST and PSI-BLAST: a new generation of protein database search programs. *Nucleic Acids Res.*, **25**, 3389–3402.
- Bandyopadhyay,S. *et al.* (2006) Systematic identification of functional orthologs based on protein network comparison. *Genome Res.*, **16**, 428–435.
- Barabasi,A.L. and Albert,R. (1999) Emergence of scaling in random networks. *Science*, **286**, 509–512.
- Boyle,E.I. *et al.* (2004) Go:termfinder open source software for accessing gene ontology information and finding significantly enriched gene ontology terms associated with a list of genes. *Bioinformatics*, **20**, 3710–3715.
- Brutlag,D.L. (2008) *Inferring Protein Function from Sequence*. Vol. 3, Chapter 30. Wiley-VCH Verlag GmbH, pp. 1087–1119.
- Camon,E. *et al.* (2004) The gene ontology annotation (GOA) database: sharing knowledge in Uniprot with gene ontology. *Nucleic Acids Res.*, **32** (Suppl. 1), D262–D266.
- Chindelevitch,L. *et al.* (2010) Local optimization for global alignment of protein interaction networks. *Pac. Symp. Biocomput.*, **10**, 123–132.
- Dezs,B. *et al.* (2011) LEMON – an open source C++ graph template library. *Electron. Notes Theor. Comput. Sci.*, **264**, 23–45. Proceedings of the Second Workshop on Generative Technologies (WGT) 2010.
- El-Kebir,M., Heringa,J. and Klau,G. (2011) Lagrangian relaxation applied to sparse global network alignment. In: Loog,M., Wessels,L., Reinders,M. and Ridder,D. (eds) *Pattern Recognition in Bioinformatics, volume 7036 of Lecture Notes in Computer Science*. Springer, Berlin Heidelberg, pp. 225–236.
- Flannick,J. *et al.* (2006) Graemlin: general and robust alignment of multiple large interaction networks. *Genome Res.*, **16**, 1169–1181.
- Flannick,J. *et al.* (2009) Automatic parameter learning for multiple local network alignment. *J. Comput. Biol.*, **16**, 1001–1022.
- Galil,Z. (1983) Efficient algorithms for finding maximal matching in graphs. In: *Proceedings of the 8th Colloquium on Trees in Algebra and Programming*. CAAP'83, Springer-Verlag, London, UK, pp. 90–113.
- Ho,Y. (2002) Systematic identification of protein complexes in *Saccharomyces cerevisiae* by mass spectrometry. *Nature*, **415**, 180–183.
- Ito,T. *et al.* (2001) A comprehensive two-hybrid analysis to explore the yeast protein interactome. *Proceedings of the National Academy of Sciences*, **98**, 4569–4574.
- Kalaev,M. *et al.* (2008) Networkblast: comparative analysis of protein networks. *Bioinformatics*, **24**, 594–596.
- Kalaev,M. *et al.* (2009) Fast and accurate alignment of multiple protein networks. *Journal of Computational Biology*, **16**, 989–999.
- Kelley,B.P. (2003) Conserved pathways within bacteria and yeast as revealed by global protein network alignment. *Proc. Natl Acad. Sci. USA*, **100**, 11394–11399.
- Kelley,B.P. *et al.* (2004) PathBLAST: a tool for alignment of protein interaction networks. *Nucleic Acids Research*, **32** (Suppl. 2), W83–W88.
- Kerrien,S. *et al.* (2012) The intact molecular interaction database in 2012. *Nucleic Acids Research*, **40**, D841–D846.
- Kirkpatrick,S. *et al.* (1983) Optimization by simulated annealing. *Science*, **220**, 671–680.
- Koyutürk,M. *et al.* (2005) Pairwise local alignment of protein interaction networks guided by models of evolution. In: Miyano,S. *et al.* (ed.) *Research in Computational Molecular Biology, volume 3500 of Lecture Notes in Computer Science*. Springer, Berlin/Heidelberg, p. 995.
- Kuchaiev,O. and Prulj,N. (2011) Integrative network alignment reveals large regions of global network similarity in yeast and human. *Bioinformatics*, **27**, 1390–1396.
- Lashkari,D.A. *et al.* (1997) Yeast microarrays for genome wide parallel genetic and gene expression analysis. *Proceedings of the National Academy of Sciences*, **94**, 13057–13062.
- Liao,C.-S. *et al.* (2009) IsoRankN: spectral methods for global alignment of multiple protein networks. *Bioinformatics*, **25**, i253–i258.
- Metropolis,N. *et al.* (1953) Equation of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*, **21**, 1087–1092.
- Morgenstern,B. (1999) Dialign 2: improvement of the segment-to-segment approach to multiple sequence alignment. *Bioinformatics*, **15**, 211–218.
- Notredame,C. *et al.* (2000) T-coffee: a novel method for fast and accurate multiple sequence alignment. *Journal of Molecular Biology*, **302**, 205–217.
- O'Brien,K.P. *et al.* (2005) Inparanoid: a comprehensive database of eukaryotic orthologs. *Nucleic Acids Research*, **33** (Suppl. 1), D476–D480.
- Pache,R.A. *et al.* (2012) NetAligner—a network alignment server to compare complexes, pathways and whole interactomes. *Nucleic Acids Research*, **40**, W157–W161.
- Park,D. *et al.* (2011) IsoBase: a database of functionally related proteins across PPI networks. *Nucleic Acids Research*, **39** (Suppl. 1), D295–D300.
- Patro,R. and Kingsford,C. (2012) Global network alignment using multiscale spectral signatures. *Bioinformatics*, **28**, 3105–3114.
- Phan,H.T.T. and Sternberg,M.J.E. (2012) Pinalog: a novel approach to align protein interaction networks implications for complex detection and function prediction. *Bioinformatics*, **28**, 1239–1245.
- Rajasekaran,S. (1990) On the convergence time of simulated annealing. Technical report. University of Pennsylvania.
- Sahraeian,S.M.E. and Yoon,B.-J. (2013) Smetana: Accurate and scalable algorithm for probabilistic alignment of large-scale biological networks. *PLoS One*, **8**, e67995.
- Schlicker,A. *et al.* (2006) A new measure for functional similarity of gene products based on gene ontology. *BMC Bioinformatics*, **7**, 302.
- Schlicker,A. *et al.* (2007) GOTax: investigating biological processes and biochemical activities along the taxonomic tree. *Genome Biol.*, **8**, R33.

- Sharan,R. (2005) Conserved patterns of protein interaction in multiple species. *Proc. Natl Acad. Sci. USA*, **102**, 1974–1979.
- Sharan,R. and Ideker,T. (2006) Modeling cellular machinery through biological network comparison. *Nat. Biotech.*, **24**, 427–433.
- Shih,Y.-K. and Parthasarathy,S. (2012) Scalable global alignment for multiple biological networks. *BMC Bioinformatics*, **13** (Suppl. 3), S11.
- Singh,R. et al. (2007) Pairwise global alignment of protein interaction networks by matching neighborhood topology. In: *Proceedings of the 11th annual international conference on Research in computational molecular biology. RECOMB'07*, Springer-Verlag, Berlin, Heidelberg, pp. 16–31.
- Singh,R. et al. (2008) Global alignment of multiple protein interaction networks with application to functional orthology detection. *Proc. Natl Acad. Sci. USA*, **105**, 12763–12768.
- Szklarczyk,D. et al. (2011) The string database in 2011: functional interaction networks of proteins, globally integrated and scored. *Nucleic Acids Res.*, **39** (Suppl. 1), D561–D568.
- Tatusov,R.L. et al. (1997) A genomic perspective on protein families. *Science*, **278**, 631–637.
- Tatusov,R.L. et al. (2000) The COG database: a tool for genome-scale analysis of protein functions and evolution. *Nucleic Acids Res.*, **28**, 33–36.
- Uniprot Consortium. (2007) The universal protein resource (UniProt). *Nucleic Acids Res.*, **35** (Suppl. 1), D193–D197.
- Waterhouse,R.M. et al. (2011) OrthoDB: the hierarchical catalog of eukaryotic orthologs in 2011. *Nucleic Acids Res.*, **39** (Suppl. 1), D283–D288.